DATABASE

PROJECT

BY
NAME: GAD NWACHUKWU OLUNGWEONWI

KALTURAL LINK: https://video.northampton.ac.uk/media/Kaltura+Capture+recording+-+May+21st+2023%2C+2A46A33+pm/1_051kcflc

Table of Contents

ASSIGNMENT REPORT **1**

1 ANALYSIS OF THE DATASET AND THE DESIGN PROCESS **3**
1.1 DATABASE NAME – **SCHOOL** 3
2 BOTH MODELS WITH GOOD EXPLANATIONS/JUSTIFICATION **4**
2.1 PURPOSE OF DATABASE 4
3 THE PROCESS OF DATABASE CREATION **5**
3.1 THE TOOL USED TO CREATE THE DATABASE IS MYSQL WORKBENCH. 5
3.2 IMPORT THE CSV DATASETS 5
3.3 THE "AWARDTITLE" TABLE 7
3.4 THE "MODULES" TABLE 7
3.5 THE "COURSES" TABLE: 8
3.6 THE PICTURE BELOW SHOWS THE ERD DIAGRAM FOR THE RELATIONAL MODEL; 8
4 CAREFULLY DESIGNED SQL AND CYPHER COMMANDS TO ANSWER THE FOLLOWING QUERIES 9
4.1 DISPLAY THOSE MODULES RUNNING IN BOTH SEMESTERS 9
4.2 DISPLAY ALL MODULES WITH AT LEAST A DESIGNATED MODULE. QUERY 10
4.3 WHICH COURSE HAS THE MOST PREREQUISITES 11
4.4 DISPLAY THOSE MODULES THAT APPEAR ON A SINGLE COURSE: QUERY: 12
4.5 DOES THE COMPUTER SCIENCE COURSE HAS UNBLANCED WORK LOAD IN BOTH SEMESTERS 13
5 ASSIGNMENT QUERIES USING GRAPH DATABASE MODEL (NEO4J) 14
5.1 N EO4J :SCHOOL 15
5.2 THE CSV FILES WERE LOADED ON THE SCHOOL DATABASE IN NO RESPECTIVE ORDER AS SHOWN BELOW IN THE NEO4J BROWSER: 17
5.3 THE FOLLOWING RELATIONSHIPS WERE CREATED AS FOLLOWS: 20
6 NEO4J QUERY 22
6 .1 DISPLAY THOSE MODULES RUNNING IN BOTH SEMESTERS: 22
6.2 DISPLAY ALL MODULES WITH ALEAST A DESIGNATED MODULE: 23
6.3 WHICH COURSE HAS MOST HAS MOST PRE_REQUISITE 24
6.4 WHICH COURSE APPEAR ON A SINGLE COURSE: 25
6.5 DOES THE COMPUTER SCIENCE COURSE HAS UNBLANCED WORK LOAD IN BOTH SEMESTERS 26
6.6 COMPARISON/REFLECTION OF RELATIONAL AND GRAPHICAL APPROACHES: 28

1 Analysis of the dataset and the design process

1.1 Database name – **school**

Tables

- ai
- ws
- cse
- cs
- bc
- cse
- modules
- courses

Columns and Data Types

- **ai, ws, cse, cs, bc, cn, se**
  - **course_code** – Integer(foreign key and related to column course_code on table courses)
  - **mod_code** – Varchar(Foreign key and related to column mod_code on table modules)
  - **requirement** – Varchar

- **modules**
  - **mod_code** – Varchar(Primary Key)
  - **old_code** – Varchar
  - **title** – Varchar
  - **credits** – Integer
  - **level** – Integer
  - **semester** – Varchar
  - **pre_requisite** – Varchar
  - **module_leader** – Varchar

- **courses**
  - **course_code** – Integer(Primary Key: AutoIncrement)
  - **award_title** – Varchar
  - **level** - Varchar

2 Both models with good explanations/justification

2.1 Purpose of Database

The purpose of the" School" database is to store and generate information about course under computer information and technology.
Courses table showcase the award given to students that study each of the 7 available courses and make us to understand each course is eligible to undergraduate students.
Modules table showcases the module attached to each of the courses and the semesters in which different modules are taken. This table also shows the leader of each module.
AI table shows all module under the Artificial Intelligence course and also their individual requirements.
BC table shows all modules under the Business Computing course and also their individual requirements.
CN table shows all modules under the Computer Networks Engineering course and also their individual requirements.
CS table shows all modules under the Computer Science course and also their individual requirements.
CSE table shows all modules under the Computer Engineering course and also their individual requirements.
SE table shows all modules under the Software Engineering course and also their individual requirements.
WS table shows all modules under the Web Development & Cyber Security course and also their individual requirements.

3 The process of database creation

3.1 The tool used to create the database is MySQL workbench.

The following are the steps followed in creating the Database.

1) **Create a Database called "school"**

   The following screenshot shows the database created using the MySQL workbench client.

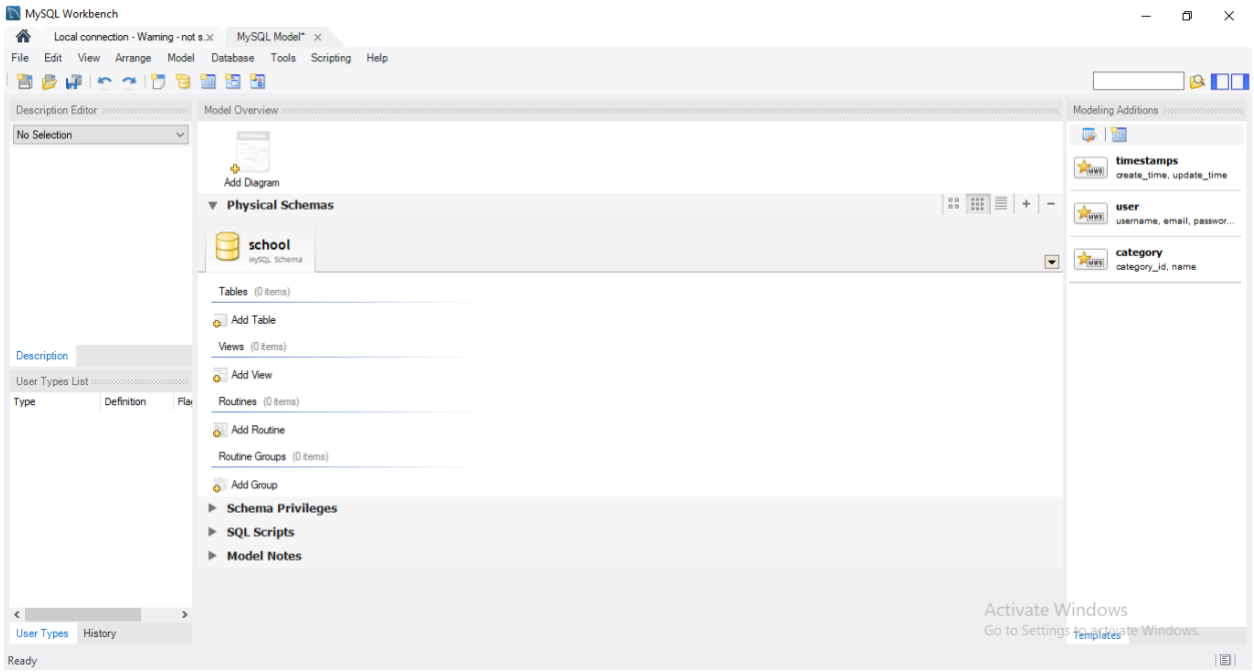

3.2 Import the CSV datasets

The following screenshots show the process of importing a csv dataset to create a table using the MySQL workbench client.

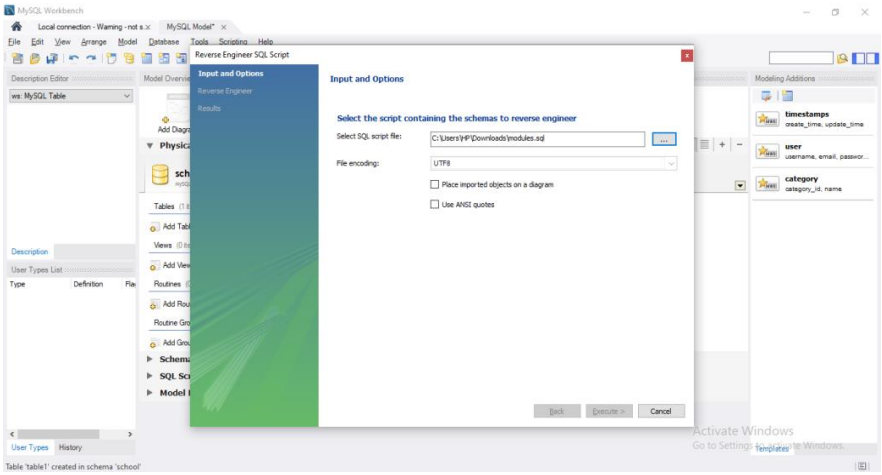

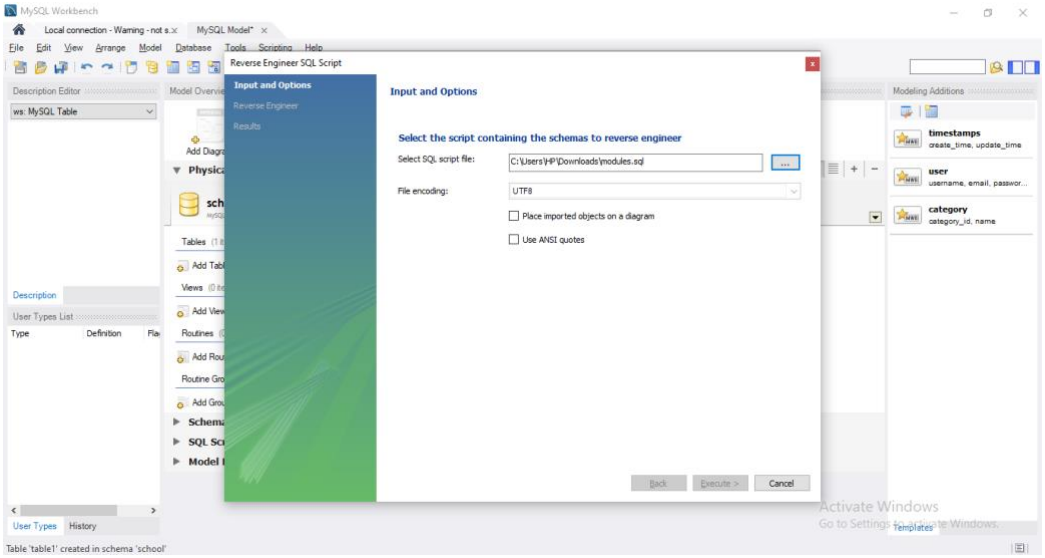

I believe that normalizing the tables is an important step in database design to improve optimization and reduce complex queries. By breaking down the original nine tables into three separate tables, I have simplified the structure and made it easier to manage the data. Here is a list of the three new tables:

The “awardtitle” Table
The “modules” Table
The “courses” Table:

By breaking the original nine tables into three separate tables, I have reduced the complexity of the database and made it easier to manage the data. This will ultimately result in better optimization and reduced complex queries when working with the data.

3.3 The “awardtitle” Table

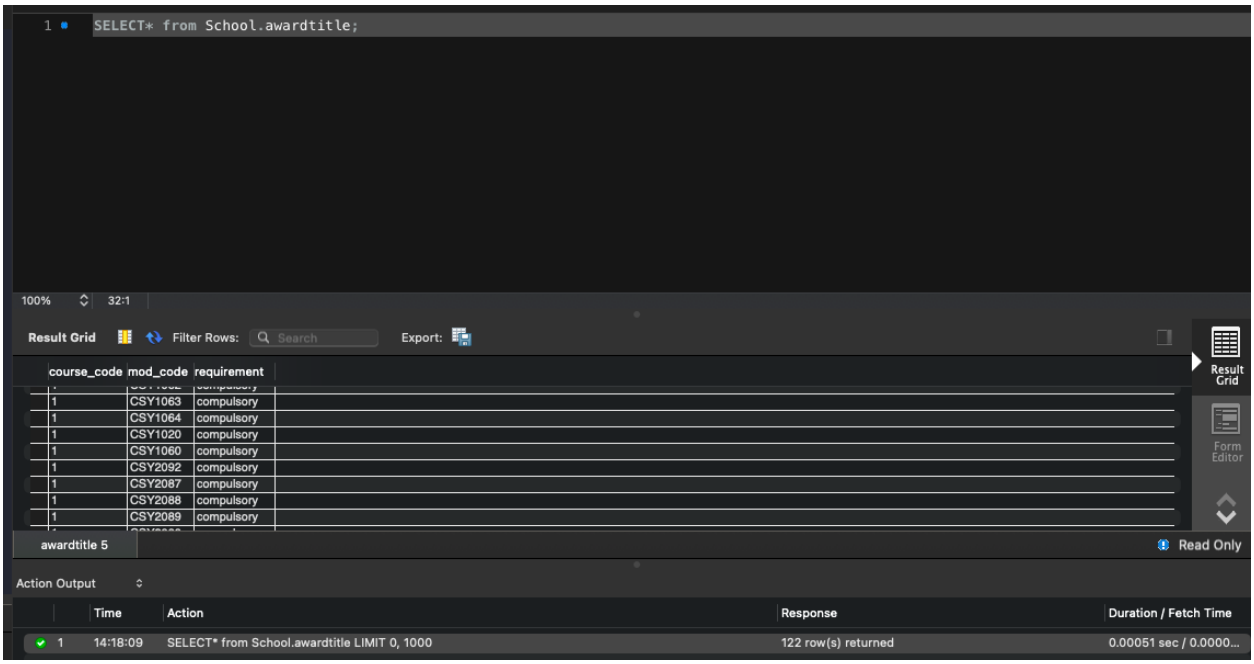

3.4 The “modules” Table

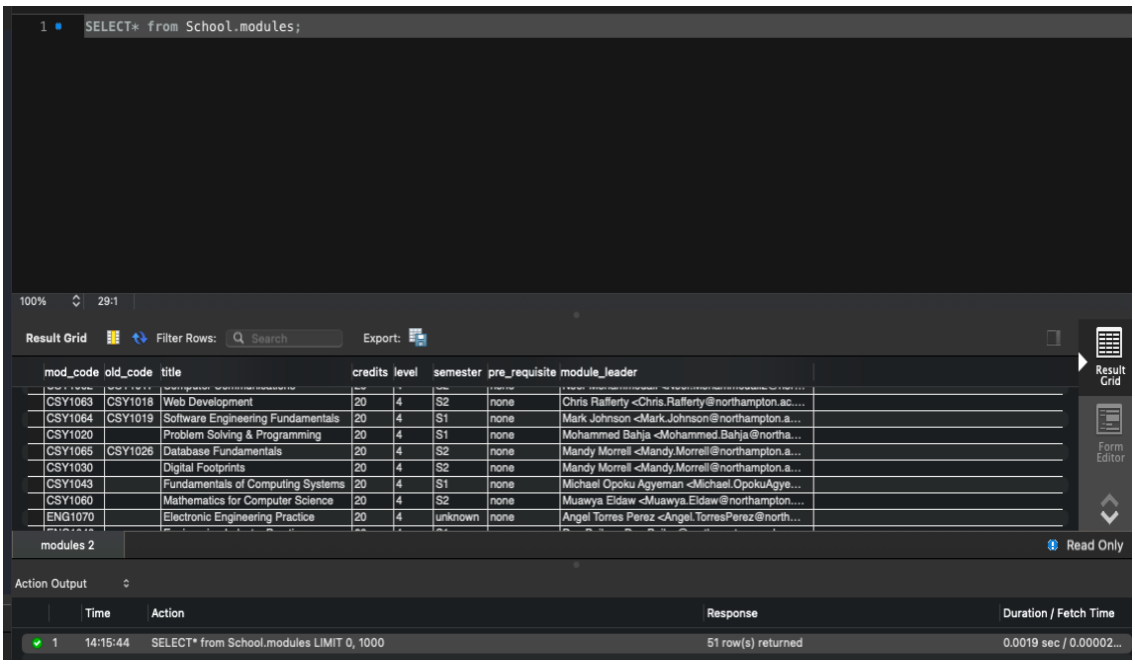

3.5 The “courses” Table:

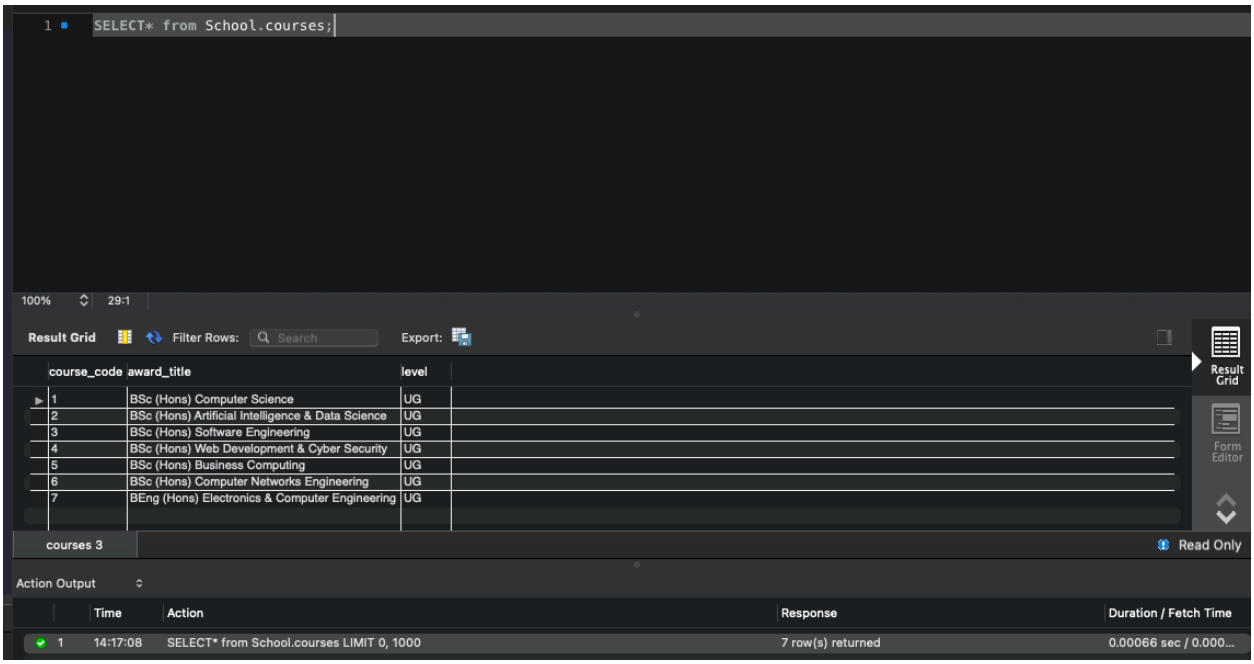

3.6 The picture below shows the ERD diagram for the relational model;

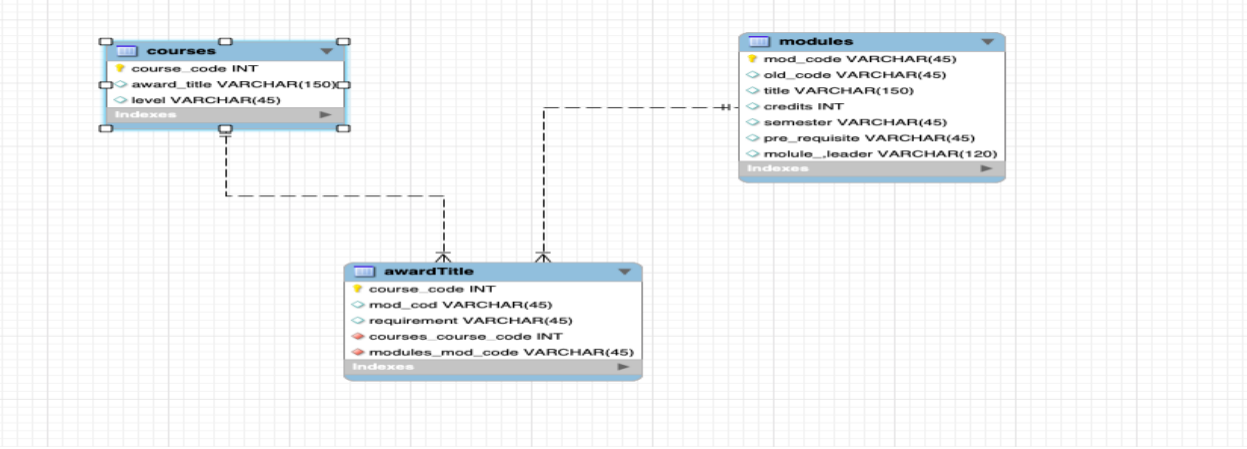

4 Carefully designed SQL and Cypher commands to answer the following queries

4.1 Display those modules running in both semesters

Query

```
SELECT mod_code, title, semester
    FROM modules
    WHERE semester = 's1 and s2';
```

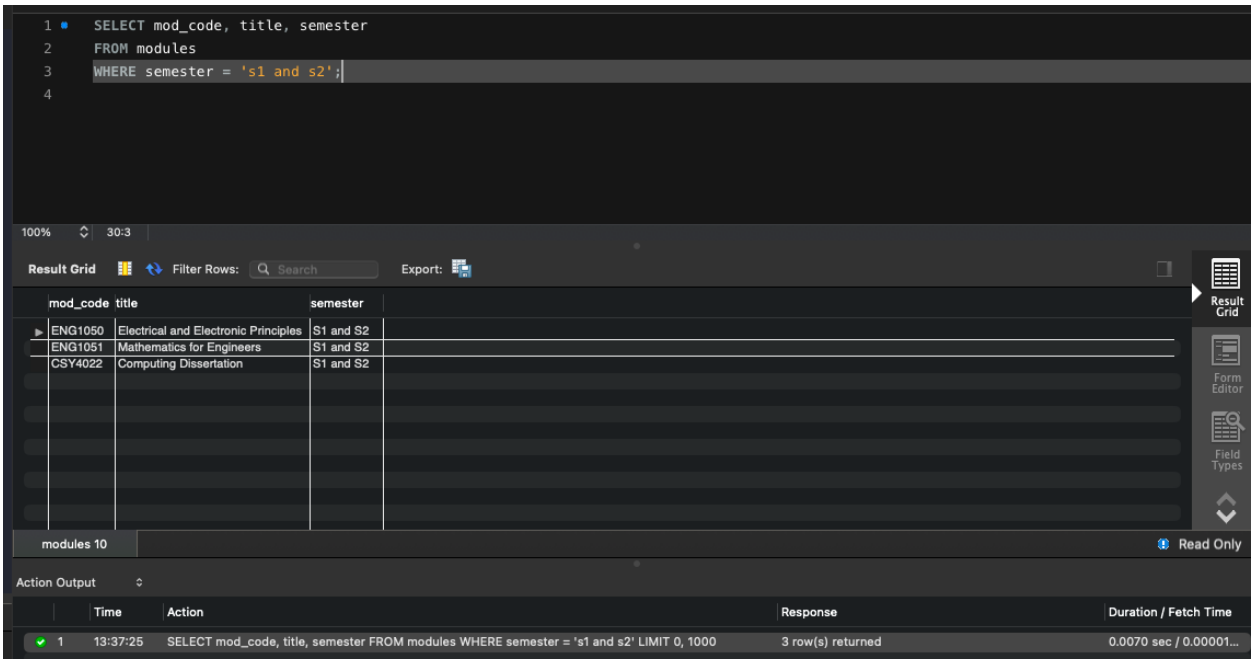

EXPLANATION:

All modules with the semester value of "s1 and s2" will have their module code, title, and semester selected by this query.

4.2 Display all modules with at least a designated module.

Query

```
SELECT course_code, mod_code, requirement
FROM awardTitle
WHERE requirement ='designated';
```

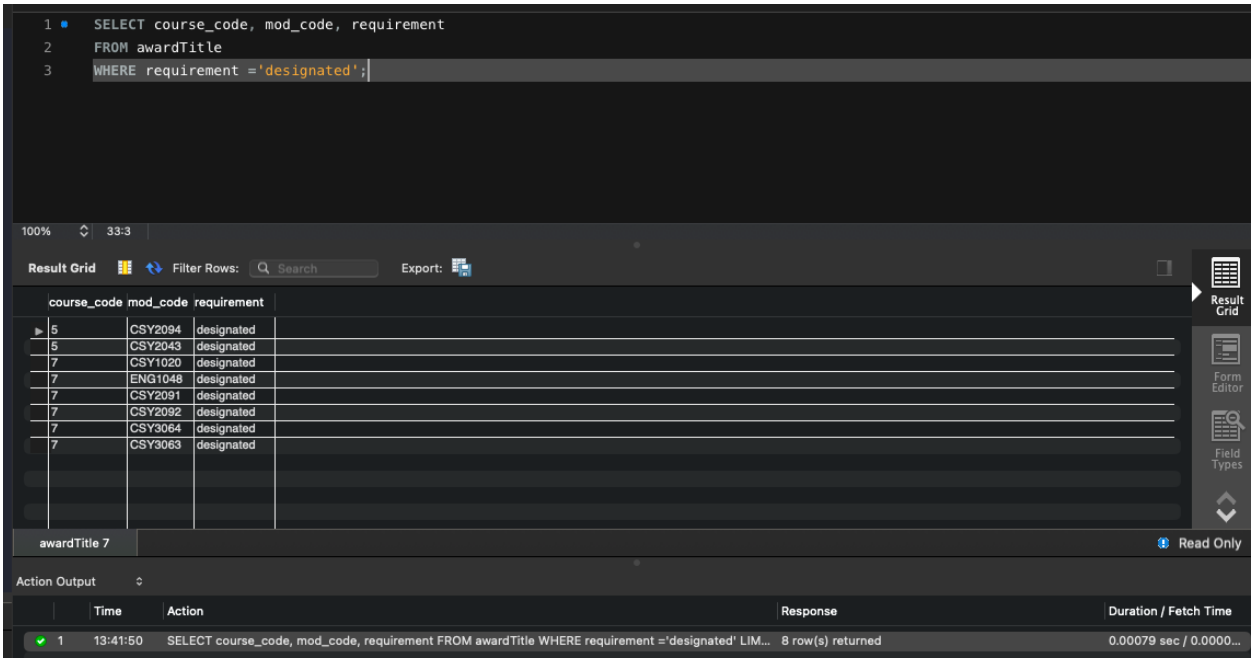

EXPLANATION:

The query retrieves the "course_code," "mod_code," and "requirement" columns from the "awardTitle" table, but only for rows where the "requirement" column is set to 'designated'.

4.3 Which course has the most prerequisites

Query:

```
SELECT c.course_code, c.award_title, COUNT(m.pre_requisite) AS prerequisite_count
FROM courses AS c
JOIN awardtitle AS cm ON c.course_code = cm.course_code
JOIN modules AS m ON cm.mod_code = m.mod_code
WHERE NOT m.pre_requisite ='none'
GROUP BY c.course_code, c.award_title
ORDER BY prerequisite_count DESC
LIMIT 1
```

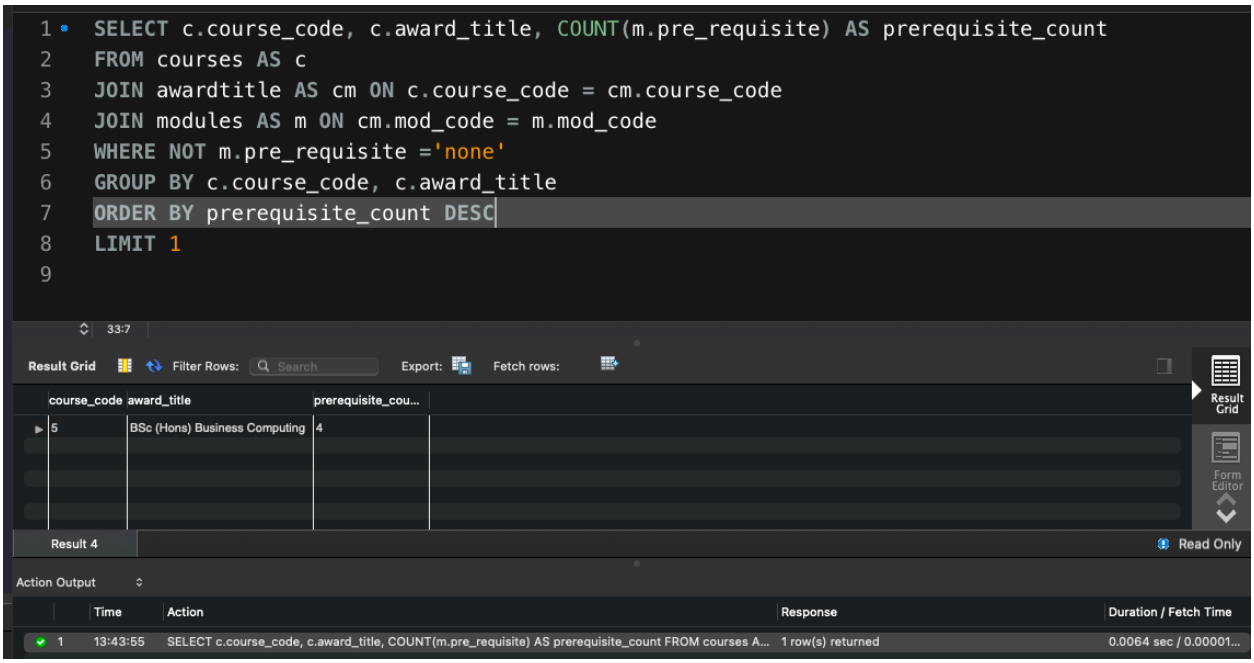

Explaination:

This query fetches details about courses such as their codes and award titles. It also calculates the number of prerequisites associated with each course. Finally, it returns the course that has the highest count of prerequisites.

4.4 Display those modules that appear on a single course:

Query:

```
SELECT m.title, m.mod_code, COUNT(a.course_code) AS COUNT
FROM courses AS c
JOIN awardtitle AS a ON c.course_code = a.course_code
JOIN modules AS m ON a.mod_code = m.mod_code
WHERE a.mod_code = m.mod_code
GROUP BY m.title, m.mod_code
HAVING COUNT = 1
ORDER BY COUNT;
```

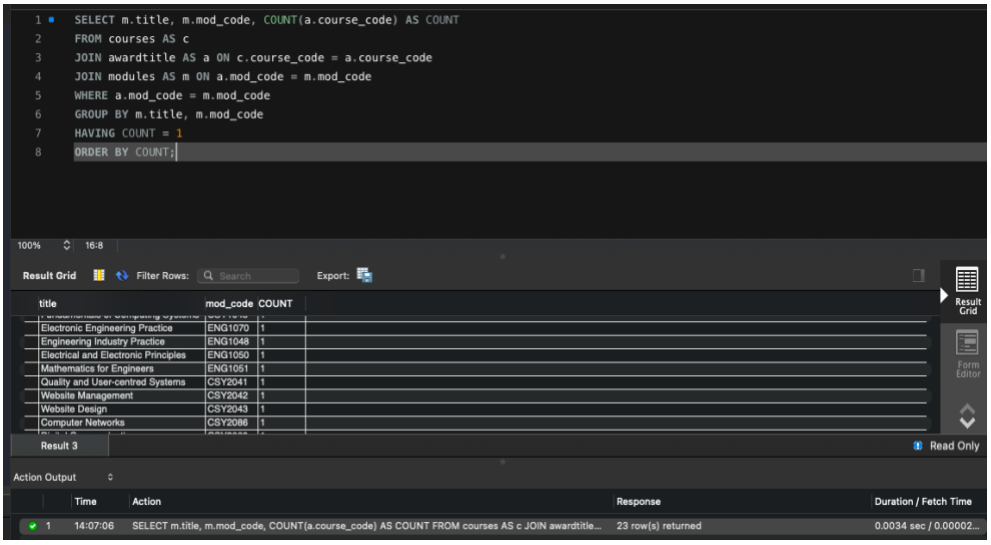

Explaination:

This query retrieves the "title" and "mod_code" columns from the "modules" table, along with the count of "course_code" from the "awardtitle" table. It performs joins between the tables, filters the results based on matching "mod_code" values, groups the results by "title" and "mod_code", and selects only the groups with a count of 1. The final result is ordered by the count in ascending order.

4.5 Does the computer science course has unblanced work load in both semesters

Query:

```
SELECT c.course_code, count(c.mod_code) AS Modules,
COUNT(CASE WHEN semester LIKE "S1" THEN "S1" END) AS S1,
COUNT(CASE WHEN semester LIKE "S2" THEN "S2" END) AS S2,
COUNT(CASE WHEN semester LIKE "S1 and S2" THEN "S1 and S2" END) AS s1_and_s2
FROM awardtitle AS c JOIN modules m ON c.mod_code = m.mod_code
JOIN courses ON courses.course_code= c.course_code
WHERE c.course_code =1
GROUP BY c.course_code
```

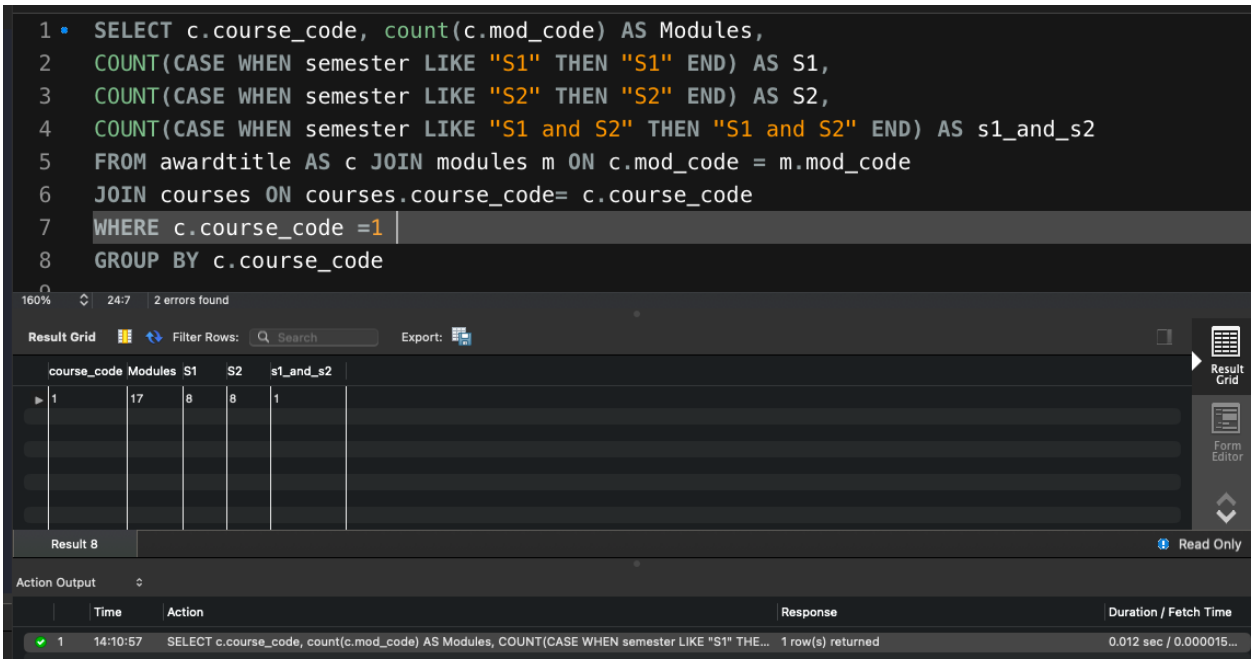

Explaination:

This query fetches the "course_code" column from the "awardtitle" table and calculates the frequency of "mod_code" for each unique "course_code". Additionally, it counts the instances of specific semesters ("S1", "S2", and "S1 and S2") by utilizing the COUNT() function and CASE statements. The query combines multiple tables, applies a filter for a particular "course_code", and groups the output based on the "course_code".

5 Assignment Queries Using Graph Database Model (Neo4j)

The dataset's analysis and the creation procedure:

The list of the given nine tables in CSV files includes modules,

I. ai,
II. cs,
III. cse,
IV. ws,
V. courses,

VI. cn,
VII. bc,
VIII. se.

5.1 N eo4j :School

The database was created and named “School” ,the data was cleaned, reformatted and moved to the import folder as shown below:

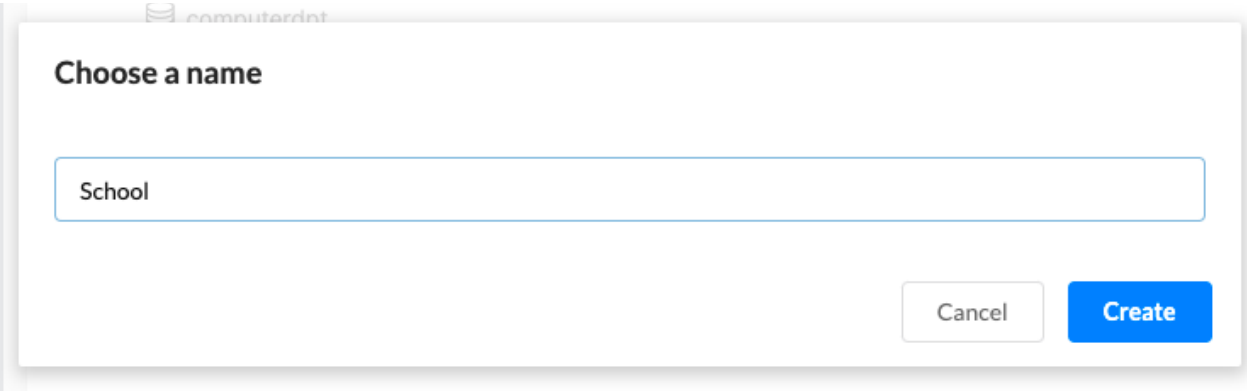

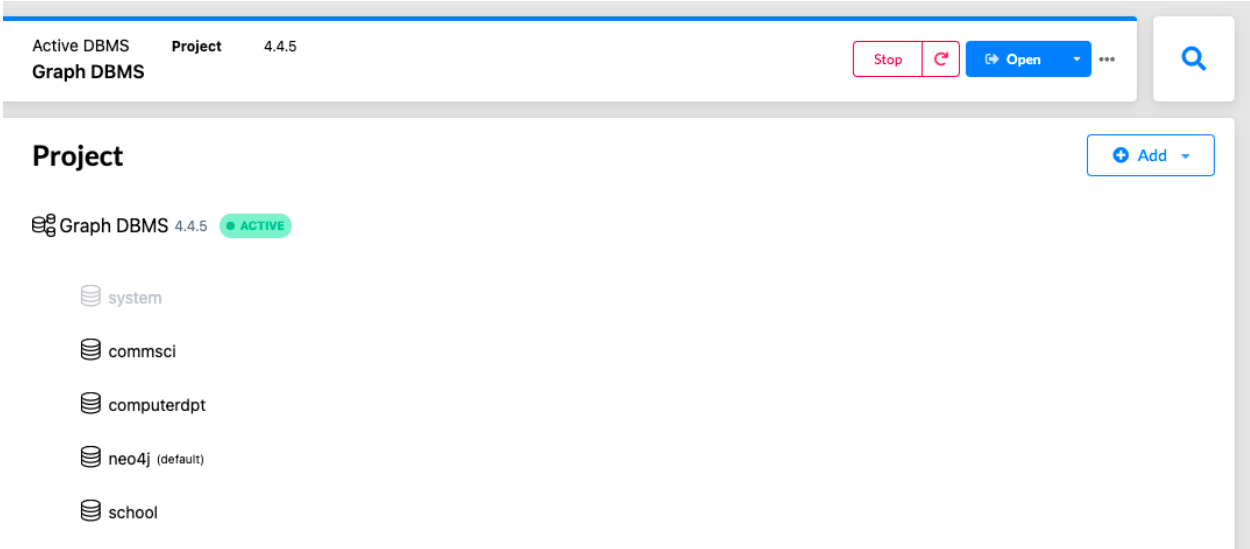

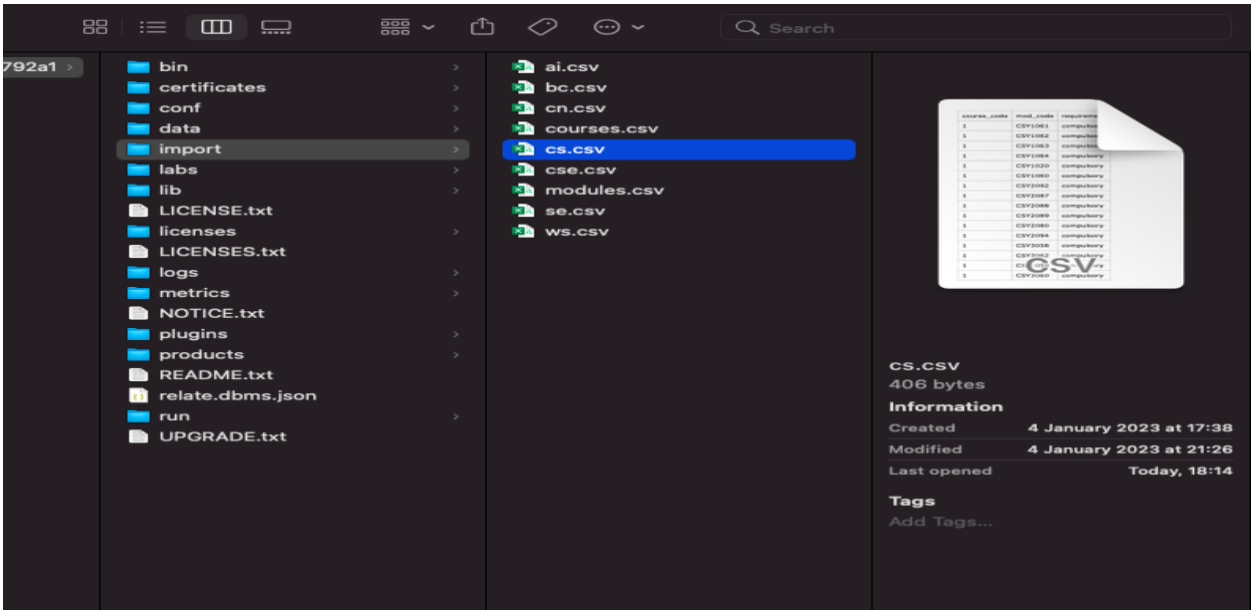
Search
792a1
bin
certificates
conf
data
import
labs
lib
LICENSE.txt
licenses
LICENSES.txt
logs
metrics
NOTICE.txt
plugins
products
README.txt
relate.dbms.json
run
UPGRADE.txt
ai.csv
bc.csv
cn.csv
courses.csv
cs.csv
cse.csv
modules.csv
se.csv
ws.csv
CSV
cs.csv
406 bytes
Information
Created 4 January 2023 at 17:38
Modified 4 January 2023 at 21:26
Last opened Today, 18:14
Tags
Add Tags...

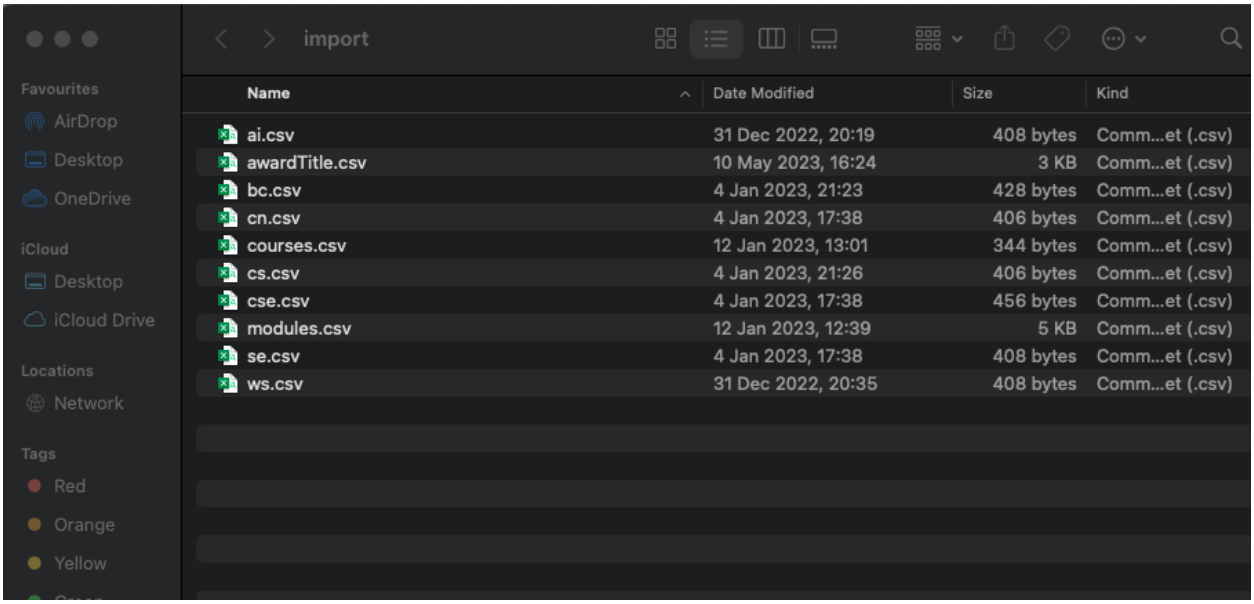
import
Favourites
AirDrop
Desktop
OneDrive
iCloud
Desktop
iCloud Drive
Locations
Network
Tags
Red
Orange
Yellow
Name
Date Modified
Size
Kind
ai.csv 31 Dec 2022, 20:19 408 bytes Comm...et (.csv)
awardTitle.csv 10 May 2023, 16:24 3 KB Comm...et (.csv)
bc.csv 4 Jan 2023, 21:23 428 bytes Comm...et (.csv)
cn.csv 4 Jan 2023, 17:38 406 bytes Comm...et (.csv)
courses.csv 12 Jan 2023, 13:01 344 bytes Comm...et (.csv)
cs.csv 4 Jan 2023, 21:26 406 bytes Comm...et (.csv)
cse.csv 4 Jan 2023, 17:38 456 bytes Comm...et (.csv)
modules.csv 12 Jan 2023, 12:39 5 KB Comm...et (.csv)
se.csv 4 Jan 2023, 17:38 408 bytes Comm...et (.csv)
ws.csv 31 Dec 2022, 20:35 408 bytes Comm...et (.csv)

5.2 The CSV files were loaded on the school database in no respective order as shown below in the Neo4j browser:

LOAD CSV WITH HEADERS FROM 'file:///modules.csv' AS line
CREATE(m:modules{mod_code:line.mod_code,old_code:line.old_code,title:line.title,credits:line.credits,level:line.level,semester:line.semester,pre_requisite:line.pre_requisite,leader:line.module_leader})
RETURN m

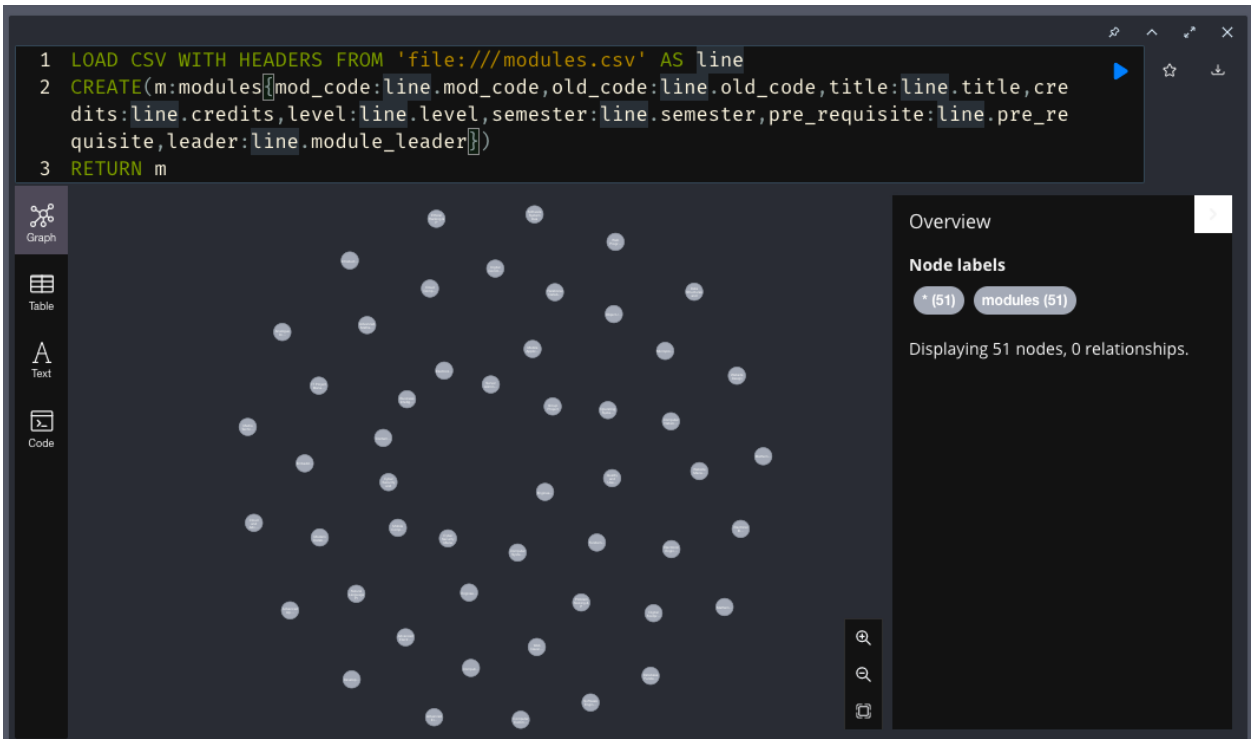

LOAD CSV WITH HEADERS FROM 'file:///awardTitle.csv' AS line
CREATE
(a:awardTitle{course_code:toInteger(line.course_code),mod_code:line.mod_code,requirement:line.requirement})
RETURN a

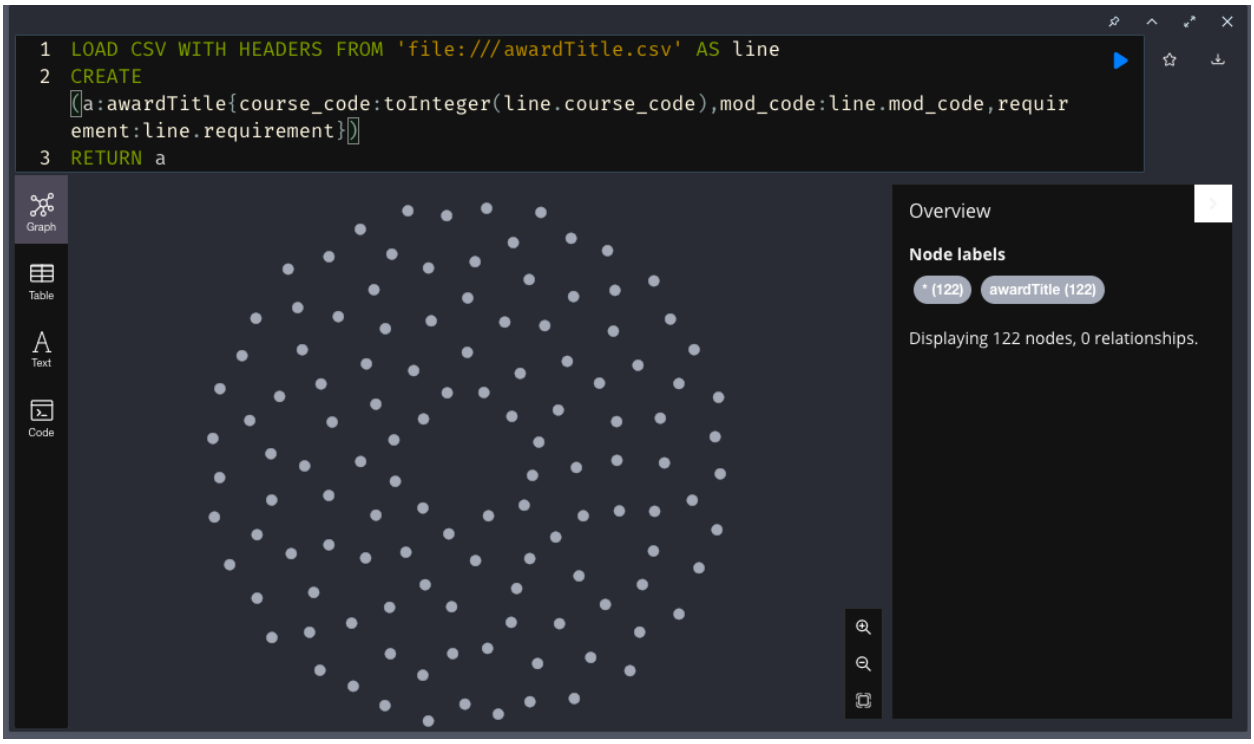

LOAD CSV WITH HEADERS FROM 'file:///courses.csv' AS line
CREATE
(c:Courses{course_code:toInteger(line.course_code),award_title:line.award_title,level:line.level})
RETURN c

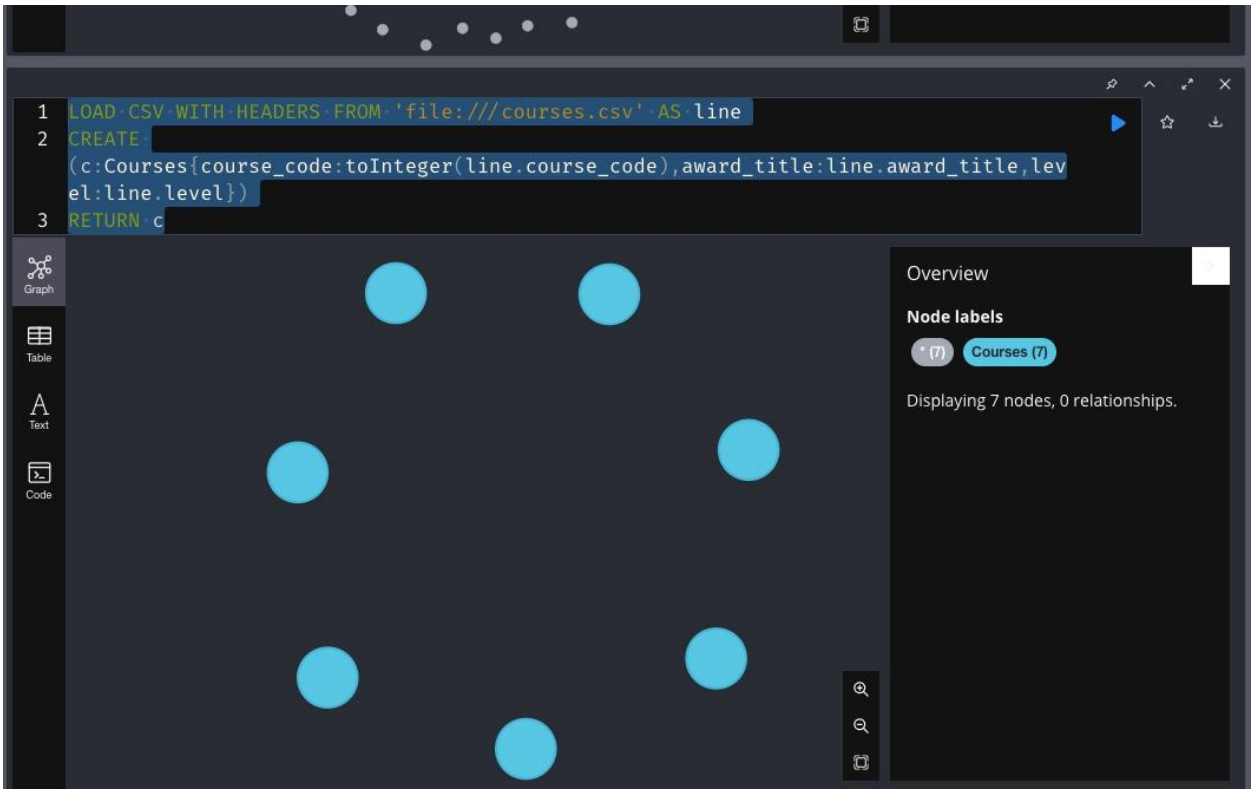

5.3 The following Relationships were created as follows:

MATCH (a:awardTitle),(m:modules)
WHERE a.mod_code = m.mod_code
MERGE (a) -[:AWARDTITLE_CONTAINS_MODULES]->(m) RETURN a,m

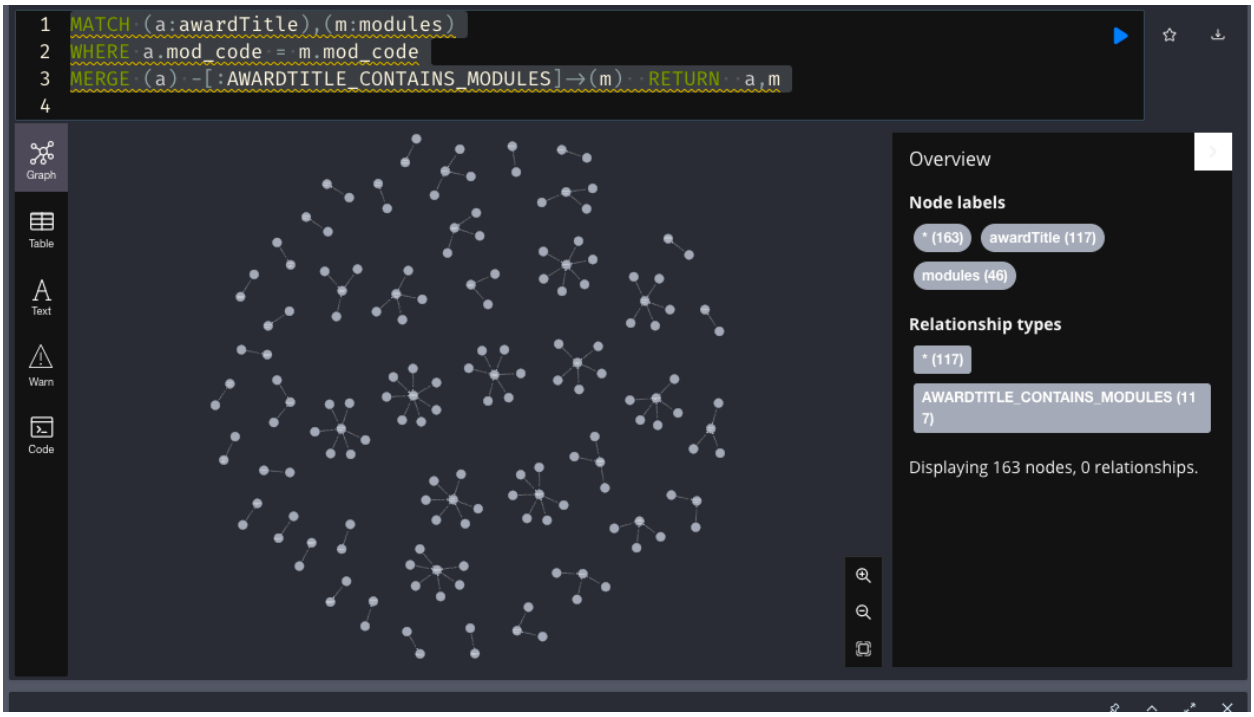

MATCH (c:Courses),(a:awardTitle)
WHERE c.course_code = a.course_code
MERGE (c) -[:COURSE_AWARDTITLE]->(a)
RETURN c, a

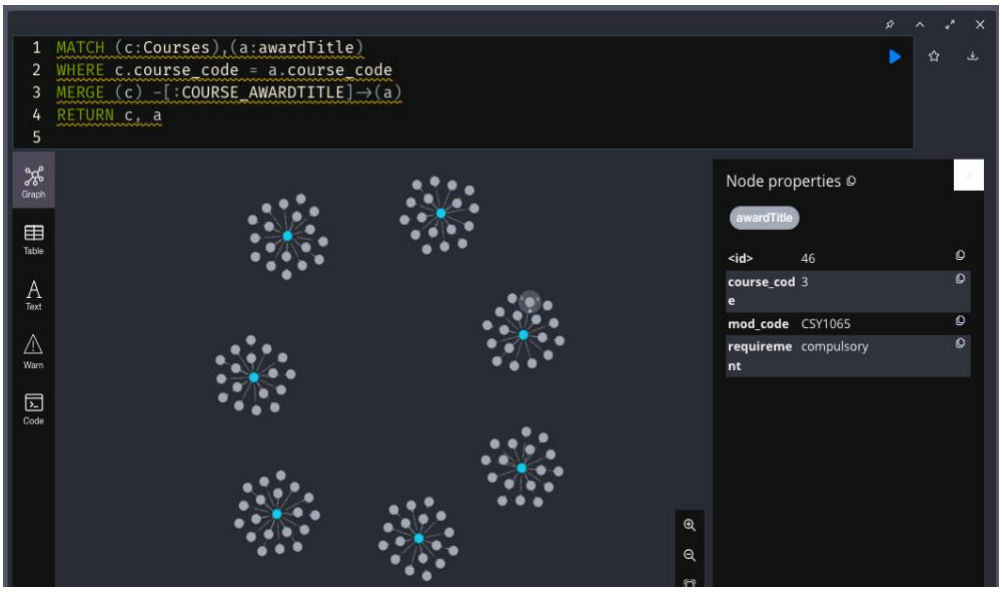

call db.schema.visualization:

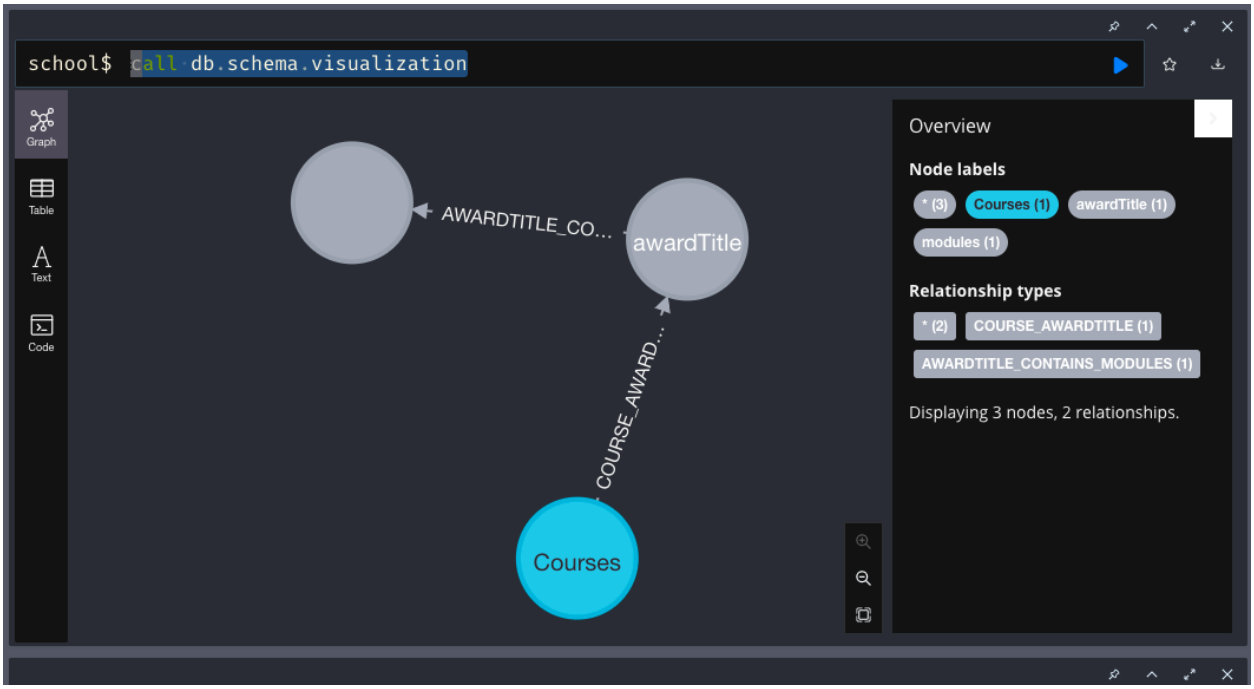

- Display those modules running in both semesters.
- Display all modules with at least a designated module.
- Which course has most pre-requisites?
- Display those modules that appear on a single course.
- Does the *Computer Science* course have an unbalanced workload in either semester?

6 Neo4j query

6 .1 Display those modules running in both semesters:

match (m:modules)
Where m.semester= 'S1 and S2'
Return m.mod_code

Query Explanation:

match (m:modules) - this line starts a match clause and defines a node pattern that matches all nodes with the modules label. The nodes are assigned the variable name m.

Where m.semester= 'S1 and S2' - this line filters the nodes based on the value of the semester property. It only returns the nodes where the semester property is equal to the string "S1 and S2".

Return m.mod_code - this line specifies the return value of the query. It returns the mod_code property of the matched nodes, in this case the modules nodes that have a semester property value of "S1 and S2".

6.2 Display all modules with aleast a designated module:

match (a:awardTitle)
Where a.requirement= 'designated'
Return a.mod_code, a.requirement

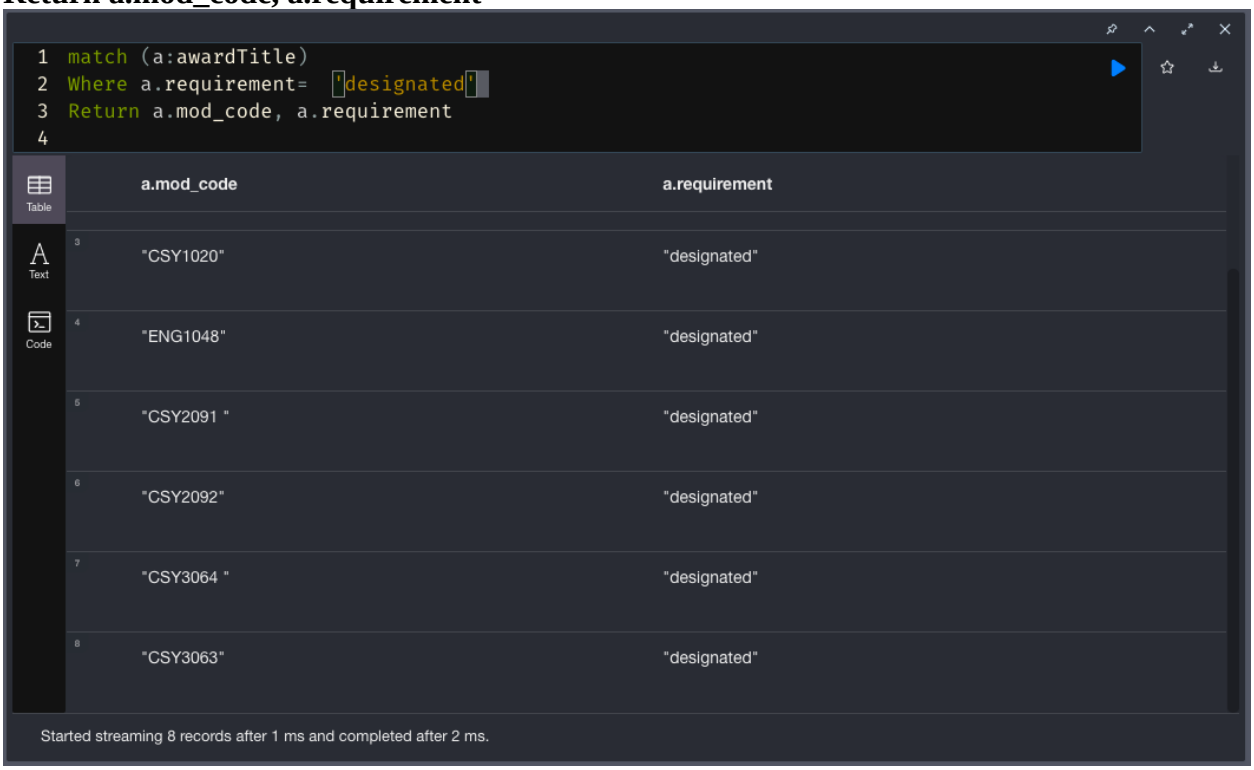

Query Explanation:

Matching all nodes with the awardTitle label is defined by the node pattern match (a:awardTitle), which begins a match clause. The variable name an is given to the nodes.

This line selects the nodes according to the value of the requirement property when a.requirement='designated'. Only nodes with the requirement attribute equal to the word "designated" are returned.

The query's return value is specified by the line Return a.mod_code, a.requirement. It returns the awardTitle nodes whose requirement property value is "designated" along with their mod_code and requirement properties.

6.3 Which course has most has most pre_requisite

```
MATCH(c:Courses)-[:COURSE_AWARDTITLE]->(a:awardTitle)-
[:AWARDTITLE_CONTAINS_MODULES]->(m:modules)
WHERE NOT m.pre_requisite = 'none'
RETURN a.course_code,c.award_title,COUNT(m.pre_requisite)
ORDER BY COUNT (m.pre_requisite) DESC
LIMIT 1
```

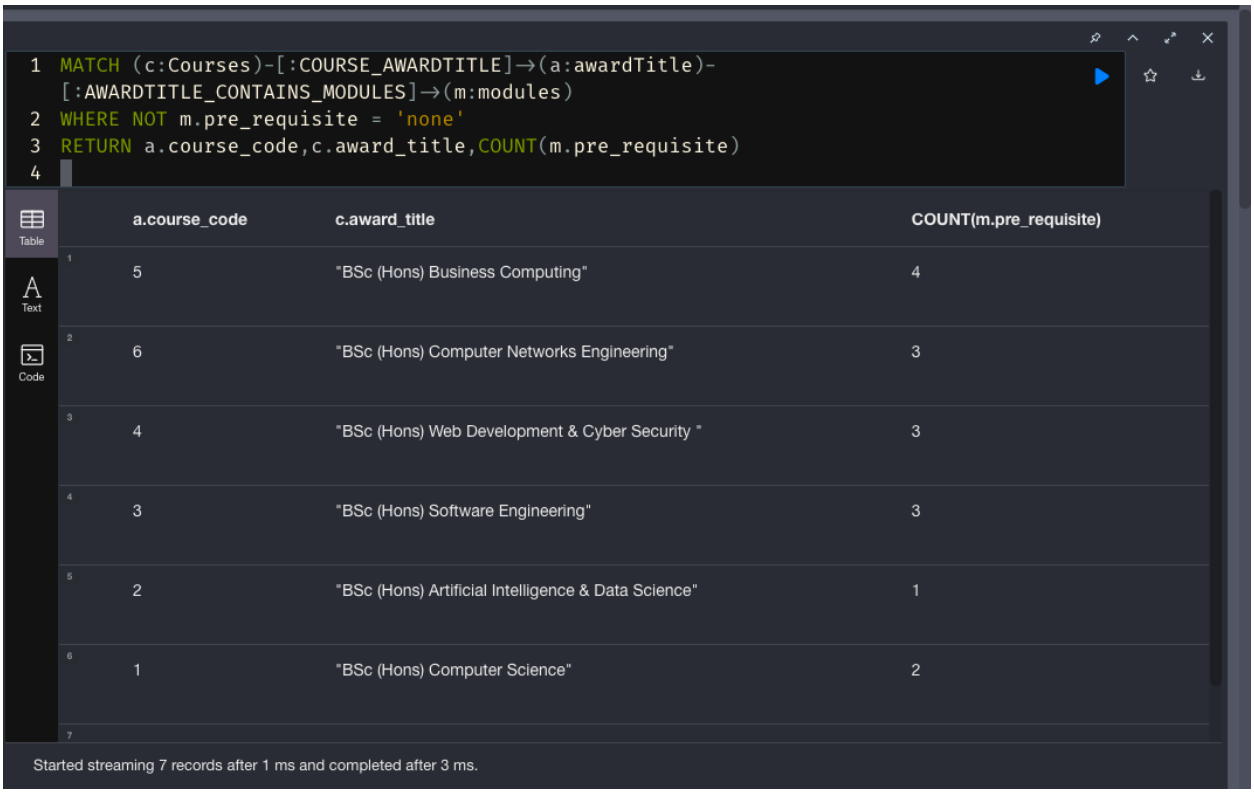

Explanation:

The query has a filter condition: WHERE NOT m.pre_requisite = 'none'. The pre_requisite property of the modules node must not be set to the value "none" in order for a path to be included.

RETURN a.course_code,c.award_title,The query's return value is indicated by the line COUNT(m.pre_requisite). It provides the award's course_code attribute back.The number of times the pre_requisite property does not equal "none" across all of the modules nodes in the matched paths, as well as the award_title and pre_requisite properties of the Courses node's award node.

M.pre_requisite ORDER BY COUNT DESC - This line uses a count of pre_requisite values to sort the results in descending order.

LIMIT 1: This line restricts the results to just the top result, which has the largest number of pre_requisite values among all the modules' nodes.

6.4 Which course appear on a single course:

MATCH (c:Courses)-[:COURSE_AWARDTITLE]->(a:awardTitle)
WITH DISTINCT (a.mod_code) as modules, collect(DISTINCT c.course_code) as courses
WHERE size(courses) = 1
Return modules,courses

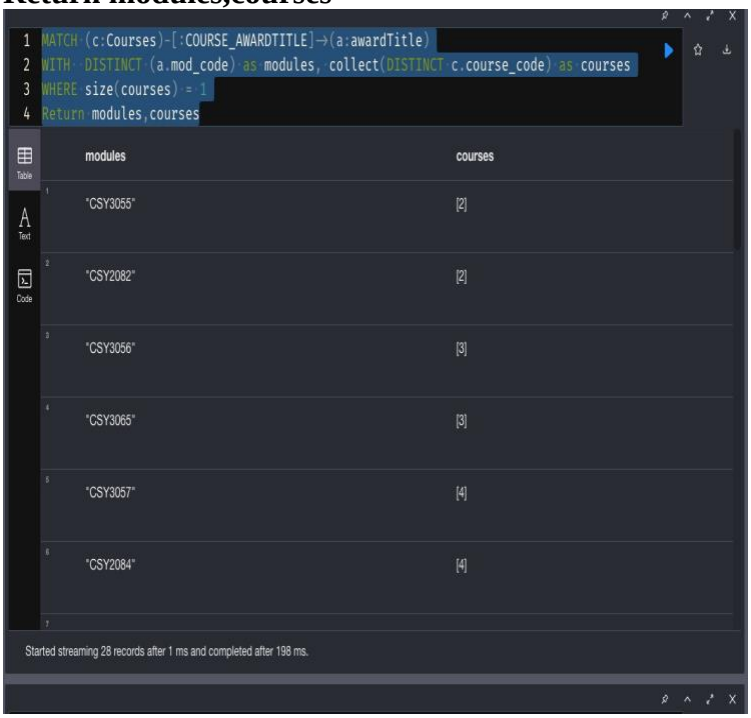

Query Explanation:

This query scans the Neo4j database for routes that begin at a node with the Courses label, travel via an outbound relationship of type COURSE_AWARDTITLE, and end at a node with the awardTitle label. The results are then organised by distinct awardTitle node mod_code values, and a list is created by distinct course_code values. Only groups with a single course linked to a certain module are included in the results after filtering. The lists of modules and courses that match the filter condition are returned at the end.

6.5 Does the computer science course has unblanced work load in both semesters

```
MATCH (c:Courses)-[rel:COURSE_AWARDTITLE]->(a)-
[r:AWARDTITLE_CONTAINS_MODULES]->(m:modules)

WITH a,m,c

WHERE c.course_code = 1

RETURN a.course_code as courses,c.award_title as
Award_title,COUNT(m.mod_code)AS ModulesCount,

COUNT(CASE m.semester WHEN "S1" THEN "S1" END) AS Semester1,

COUNT(CASE m.semester WHEN "S2" THEN "S2" END) AS Semester2,

COUNT(CASE m.semester WHEN "S1 and S2" THEN "S1 and S2" END) AS
Semester1_AND_2

ORDER BY Semester1,Semester2

ORDER BY Semester1,Semester2
```

```
MATCH (c:Courses)-[rel:COURSE_AWARDTITLE]→(a)-[r:AWARDTITLE_CONTAINS_MODULES]→
(m:modules)
WITH a,m,c
WHERE c.course_code = 1
RETURN a.course_code as courses,c.award_title as Award_title,COUNT(m.mod_code)AS
ModulesCount,
COUNT(CASE m.semester WHEN "S1" THEN "S1" END) AS Semester1,
COUNT(CASE m.semester WHEN "S2" THEN "S2" END) AS Semester2,
COUNT(CASE m.semester WHEN "S1 and S2" THEN "S1 and S2" END) AS Semester1_AND_2
ORDER BY Semester1,Semester2
```

Table | Text | Code

	courses	Award_title	ModulesCount	Semester1	Semester2	Semester1_AND_2
1	1	"BSc (Hons) Computer Science"	17	8	8	1

Started streaming 1 records after 360 ms and completed after 363 ms.

Query Explanation:

This search finds all the modules nodes associated with a node with a course_code of 1, counts the number of modules in each semester type, and returns the results. Following that, the query returns the course_code of the n node as computer_science, the total number of mod_code values for the modules nodes as ModulesCount, and the number of semester values for the modules nodes for each semester type. The query then arranges the outcomes according to the Semester1 and Semester2 columns in ascending order.

6.6 Comparison/Reflection of relational and graphical approaches:

The relational database is primarily used for small organisational databases because it is structured in a table and is simple to query. The relational database captures all properties or features entities and makes them unique, but when the data is large and involves several relationships, the relational database cannot capture this efficiently, particularly how one entity relates to another. As a result, non-relational or graph databases are required. This can be observed in databases used by firms such since Facebook, Amazon, and others, as emerging technologies such as Neo4j are particularly efficient with data connections. A look at my assignment above demonstrates how the graph Neo4j was able to capture the relationship latency in SQL queries.